

23.01.2017

18.18 – Hakkan linuxiga jälle tegelema.

Paneme folderitd kõik nähtavaks. Kustutasime ja shrinkisime foldereid. Tahaks suure restardi teha. Uurisin, et võimalik linux mittejärjesikustela partitsioonidele ühtlaselt installid .

Sisuliselt 2 võimalus

- 1) Kaotame kuidagi partitisooonid
- 2) Installime Linuxi mitmela partitsioonile

01.02.2017

1. Kasutasin sellist programmin nagu EaseUS partitsioneerimiseks
2. Installisin Ubuntu USB-le <http://www.howtogeek.com/howto/linux/create-a-bootable-ubuntu-usb-flash-drive-the-easy-way/>
3. <https://courses.cs.ut.ee/2016/os/fall/Main/Praktikum2> Hakkan selle järgi installima ubuntu
a. Enne tutvun kindlamalt partitsioonidega

Partitsioonid

<https://www.bleepingcomputer.com/tutorials/understanding-hard-disk-partitions/>
- hea tutorial

In the current IBM PC architecture, there is a partition table in the drive's Master Boot Record (section of the hard drive that contains the commands necessary to start the operating system), This partition table is then **further split into 4 partition table entries**, with each entries corresponding to a partition. This partition table is then further split into 4 partition table entries, with each entries corresponding to a partition. To overcome this restriction, system developers decided to add a new type of partition called the **extended partition**. By replacing one of the four primary partitions with an extended partition, you can then make an additional **24 logical partitions** within the extended one.

Partition Table

Primary Partition #1

Primary Partition #2

Primary Partition #3

Primary Partition #4 (Extended
Partition)

Logical Partition #1

Logical Partition #1

Each hard drive also has one of its possible 4 partitions flagged as an **active partition**. The active partition is a special flag assigned to only one partition on a hard drive that the Master Boot Record (MBR) uses to boot your computer into an operating system

As only one partition may be set as the active partition, you may be wondering how people can have multiple operating systems installed on different partitions, and yet still be able to use them all. This is accomplished by installing a boot loader in the active partition. When the computer starts, it will read the MBR and determine the partition that is flagged as active. This partition is the one that contains the boot loader

When the operating system boots off of this partition the boot loader will start and allow you to choose which operating systems you would like to boot from

16:48 Mounting a fileSystem in Linux

<https://www.bleepingcomputer.com/tutorials/introduction-to-mounting-filesystems-in-linux/>

In order to access a filesystem in Linux you first need to mount it. Mounting a filesystem simply means making the particular filesystem accessible at a certain point in the Linux directory tree. When mounting a filesystem it does not matter if the filesystem is a hard disk partition, CD-ROM, floppy, or USB storage device.

You simply need to know the device name associated with the particular storage device and a directory you would like to mount it to.

. Having the ability to mount a new storage device at any point in the directory is very advantageous. For example, lets say that you have a web site stored in /usr/local/website. The web site has become very popular and you are running out of space on your 36 GB hard drive. You can simply go out and purchase a new 73 GB hard drive, install it in the computer, and then mount that entire drive as /usr/local/. Now your /usr/local mount point has a total hard drive space of 73 GB, and you can free up the old hard drive by copying everything from the old /usr/local to the new one. As you can see, adding more hard drive space to a computer, while still keeping the same exact directory structure, is now very easy.

16:54 Nüüd edasi tulevad käsud, seega vaatan, et saaksin usb pealt bootida.

<http://www.isunshare.com/windows-password/four-methods-to-access-uefi-bios-setup.html>

17.09 Sain käima UBUNTU TRY osa. Proovin nüüd ubuntu partitsionides aru saada.

<https://help.ubuntu.com/community/PartitioningSchemes> MINGI GPT jama teksti alguses. Uurin mis schemen mul on. **Partition Scheme** keyword

<http://www.howtogeek.com/howto/35676/how-to-choose-a-partition-scheme-for-your-linux-pc/>

17.24 Installin Ubuntut eelnevalt mainitud tutoriali järgi.

Seal õpetatakse, et kui insallima hakkan siis

Boot jaoks 32 Swap jaoks 4 ja Home ülejäänud

17.29 – 15 mintsa paus.

02.02.2017

9.44

<http://linuxbsdos.com/2015/10/30/gpt-and-mbr-manual-disk-partitioning-guide-for-ubuntu-15-10/>

<http://linuxbsdos.com/2014/11/08/a-beginners-guide-to-disks-and-disk-partitions-in-linux/>

Newer computers come with a replacement firmware for the old BIOS system called UEFI (Unified Extensible Firmware interface), and GPT is a part of the UEFI standard. If you bought a Windows 8 computer, it's most certainly installed on a GPT partitioning scheme. If you're already running a Linux distribution or attempting to install one on a recent Windows 7/8 computer, the easiest method of finding out what partitioning standard is used is to launch a shell terminal and type **sudo fdisk -l** (leave out the *sudo* if the distribution does not use it).

sudo fdisk -l

Mul SDA -l gpt

File Systems: Before a disk partition can be used to store data, it must first be formatted. The formatting process includes stamping it with a file system.

10.27 – Peaks natuke tegelema mõistetega ja läbi mõtlemiseg..

GPT MBR

Partitsioonid linuxis

LVM linuxis

Linuxi installimine on võimalik ilma LVM-ita Võimalik koos LVM-iga.

Linuxil on mount pointid . Alguseks oluline installida boot, home ja swap.

Failisüsteem kas ext 4 või bdfs

Põhiküsimused

1. Kas kasutada LVM

Mis on grub ja mis on bootloadeR This partition is the one that contains the boot loader

When the operating system boots off of this partition the boot loader will start and allow you to choose which operating systems you would like to boot from

- 2.
- 3.

10.52

1. 1. Kas kasutada LVM – EI
2. Mis on EFI, KAS MUL ON VAJA BOOTLOADERI JAOKS MIDAGI ERALDI
 - a. Kas on efit vaja And if you're attempting to set up a dual-boot system between Windows 10 and Ubuntu 15.10 on a single hard drive, you'll be creating just three partitions – each mounted at / (root), /home and swap. That means creating an EFI System Partition (ESP), a root, /home and swap partitions. The /home partition is optional but recommended. => Seega minul ei ole vaja

<http://www.tecmint.com/ubuntu-16-04-installation-guide/#C3>

<http://linuxbsdos.com/2015/10/30/gpt-and-mbr-manual-disk-partitioning-guide-for-ubuntu-15-10/>

Boot loader i

<http://askubuntu.com/questions/326662/which-partition-to-select-as-device-for-boot-loader-in-an-efi-mode-install>

<https://ubuntuforums.org/showthread.php?t=2309806>

Ubutntu installitud valisin efi bootloaderiks. Praefu peab bioses manuaalselt järjekorda vahetama.

04.02.2017

Apt-get update

Apt-get install ssh

11.09

SSH -ga tegelen .

Praegu ühendatud läbi kahe arvuti. Kuidas faile transportida?

11:33 Avasin tutoriali 1

11.57 -Definition of hypervisor.

On 2-te tüüpi.

Type-1, native or bare-metal hypervisors

These hypervisors run directly on the host's hardware to control the hardware and to manage guest operating systems. For this reason, they are sometimes called bare metal hypervisors. The first hypervisors, which IBM developed in the 1960s, were native hypervisors.^[4] These included the test software SIMMON and the CP/CMS operating system (the predecessor of IBM's z/VM). Modern equivalents include Xen, Oracle VM Server for SPARC, Oracle VM Server for x86, the Citrix XenServer, Microsoft Hyper-V and VMware ESX/ESXi.

Type-2 or hosted hypervisors

These hypervisors run on a conventional operating system just as other computer programs do. A guest operating system runs as a process on the host. Type-2 hypervisors abstract guest operating systems from the host operating system. VMware Workstation, VMware Player, VirtualBox, Parallels Desktop for Mac and QEMU are examples of type-2 hypervisors.

<https://en.wikipedia.org/wiki/Hypervisor>

<https://en.wikipedia.org/wiki/Xen>

11.37 – Vahetasin parooli. Esimene ühendus tehtud

Tahaks rohkem teada kuidas ühendused toimivad.

1) Try to connect in LAN with gateguard

Exactly like in tutorial with putty .

1

Gatekeeper.cs.hku.hk Port: 22

anamm

AedrMDwH

2

ssh student@202.45.128.135

3

ssh student@student13 tundis ilusti ära

2) Proovin ühenda võrgus ilma gateguardita. - El

202.45.128.135

student@202.45.128.135

3) Proovin ühendada võrgust väljas ilma gateguardita, ilma vpn-ita - El

4) Proovin ühendada võrgust väljas gateguardita ja vpn-iga - SOBIB

ssh student@202.45.128.135

⇒ Ilmselt on vaja olla CS department vpn-is

5) Proovin ühendada võrgust väljas gateguardi ja vpn-iga SOBI

6) Proovin ühendada VÕRGUST VÄLJAS ilma vpnita gateguardiga SOBIB

Täpselt sama, kui kõik muu

13.28 FAILIDE LIIGUTAMINE TOIMUB VAID LÄBI VPN-I SOBIB.

16.27

See teine vahelüli võrgus

- NAT https://en.wikipedia.org/wiki/Network_address_translation
<http://whatismyipaddress.com/nat>

Network address translation (NAT) is a method of remapping one IP address space into another by modifying network address information in Internet Protocol (IP) datagram packet headers coc server

- DHCP server [https://technet.microsoft.com/en-us/library/dd145320\(v=ws.10\).aspx](https://technet.microsoft.com/en-us/library/dd145320(v=ws.10).aspx) Automatic ip address mappin
- TFTP server FILE RELATED

17.04

Part I Specification

sudo lshw -short

lshw(Hardware Lister): a small tool to provide detailed information on the hardware configuration of the machine

H/W path	Device	Class	Description
----------	--------	-------	-------------

=====

	system		HP Compaq Elite 8300 SFF (QV996AV)
/0	bus	3397	
/0/0	memory		64KiB BIOS
/0/4	memory		256KiB L1 cache
/0/5	memory		1MiB L2 cache
/0/6	memory		6MiB L3 cache
/0/7	memory		16GiB System Memory
/0/7/0	memory		DIMM [empty]

/0/7/1 memory 8GiB DIMM DDR3 Synchronous 1600 MHz (0.6 ns)

<http://www.pcmag.com/article2/0,2817,2400801,00.asp>

/0/7/2 memory DIMM [empty]

/0/7/3 memory 8GiB DIMM DDR3 Synchronous 1600 MHz (0.6 ns)

/0/10 processor Intel(R) Core(TM) i5-3570 CPU @ 3.40GHz

<https://arstechnica.com/gadgets/2016/02/pentium-core-i5-core-i7-making-sense-of-intels-convoluted-cpu-lineup/>

These are all quad-core CPUs without **Hyperthreading**,

/0/100 bridge Xeon E3-1200 v2/3rd Gen Core processor DRAM Controller

/0/100/2 display Xeon E3-1200 v2/3rd Gen Core processor Graphics Controller

/0/100/14 bus 7 Series/C210 Series Chipset Family USB xHCI Host Controller

/0/100/19 eth0 network 82579LM Gigabit Network Connection

/0/100/1a bus 7 Series/C210 Series Chipset Family USB Enhanced Host Controller #2

/0/100/1b multimedia 7 Series/C210 Series Chipset Family High Definition Audio Controller

/0/100/1d bus 7 Series/C210 Series Chipset Family USB Enhanced Host Controller #1

/0/100/1e bridge 82801 PCI Bridge

/0/100/1f bridge Q77 Express Chipset LPC Controller

/0/100/1f.2 storage 7 Series/C210 Series Chipset Family 6-port SATA Controller [AHCI mode]

```

/0/100/1f.3      bus      7 Series/C210 Series Chipset Family SMBus Controller
/0/1      scsi0      storage
/0/1/0.0.0    /dev/sda  disk      500GB ST500DM002-1BD14
/0/1/0.0.0/1  /dev/sda1  volume    449GiB EXT4 volume
/0/1/0.0.0/2  /dev/sda2  volume    15GiB Extended partition
/0/1/0.0.0/2/5 /dev/sda5  volume    15GiB Linux swap / Solaris partition
/0/2      scsi2      storage
      /0/2/0.0.0    /dev/cdrom disk      DVD-RAM GH80N

```

QUESTION – DOES IT HAVE EFI OR MBR – Probably mbr, because we use an extended partition

```
cat /proc/cpuinfo
```

```
cat /proc/cpuinfo | grep processor | wc -l
```

You need to use Proc (/proc) file system provides information about CPU and their speed which is a pseudo-filesystem <https://www.cyberciti.biz/faq/linux-display-cpu-information-number-of-cpus-and-their-speed/>

lcpuhy it shows multiple cpu-s in this file

<http://unix.stackexchange.com/questions/146051/number-of-processors-in-proc-cpuinfo>

The words “CPU”, “processor” and “core” are used in somewhat confusing ways. They refer to the processor architecture.

A core is the smallest independent unit that implements a general-purpose processor;

A processor is an assemblage of cores (on some ARM systems, a processor is an assemblage of clusters which themselves are assemblages of cores).

A chip can contain one or more processors (x86 chips contain a single processor, in this sense of the word *processor*).

Hyperthreading means that some parts of a core are duplicated. A core with hyperthreading is sometimes presented as an assemblage of two “virtual cores”

In software manuals, terms *CPU* and *processor* are used to mean any one piece of hardware that executes program code. In hardware terms, this means one core, or one virtual core with hyperthreading.

Hardware configuration in STUDENT 13

CPU model: Intel(R) Core(TM) i5-3570 CPU @ 3.40GHz

RAM size: memory 16GiB System Memory

of cores : 4 No hyperthreading

XEN installation

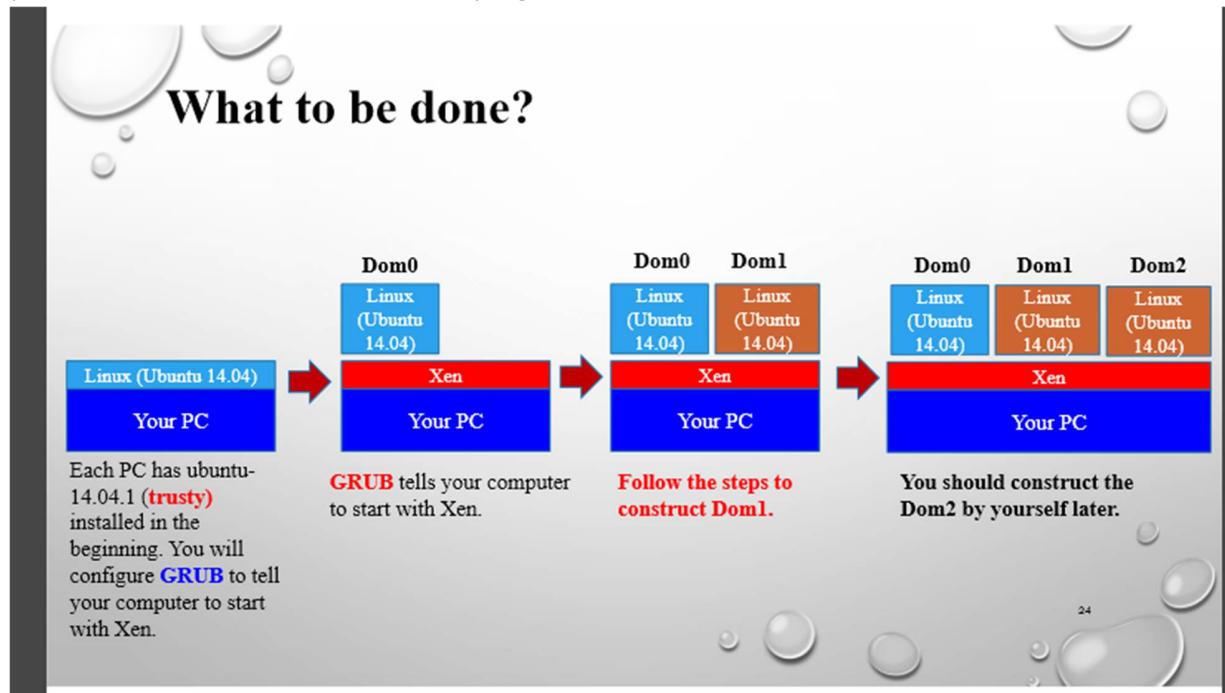
18:38

1. Started with the slides
2. Started with the video in moodle
3. Went on to tutorial

https://wiki.xenproject.org/wiki/Xen_Project_Beginners_Guide

18.44

punane nool siin slaidil tähistab järgmisele sammule liikumist.



05.02.2015

Network switch https://en.wikipedia.org/wiki/Network_switch Does stuff in OSI layer 2

<https://www.quora.com/What-is-the-difference-between-HTTP-protocol-and-TCP-protocol>

19.05

sudo apt-get update

sudo apt-get upgrade

xen

sudo apt-get install xen-hypervisor-amd64

vim

sudo apt-get install vim

READING ABOUT BOOTLOADER

https://en.wikipedia.org/wiki/GNU_GRUB

<http://www.howtogeek.com/196655/how-to-configure-the-grub2-boot-loaders-settings/>

Distribution info

lsb_release -a

Distributor ID: Ubuntu

Description: Ubuntu 14.04.5 LTS

Release: 14.04

Codename: trusty

GRUB settings

sudo vim /etc/default/grub

https://www.gnu.org/software/grub/manual/html_node/Simple-configuration.html

GRUB_CMDLINE_LINUX'

Command-line arguments to add to menu entries for the Linux kernel.

Do a copy first

sudo cp /etc/default/grub cp_grubfile

sudo

<http://www.openvim.com/>

Väike vim tuts

sudo update-grub

Network configuration

virtual switch – like a physical switch, but virtual - Dom0(student10) that takes packets from the virtual machines (Dom1 & Dom2) and forwards them onto the

physical network so they can see the Internet and other machines on your network.

We use linux bridge for this

https://wiki.xenproject.org/wiki/Xen_Project_Beginners_Guide

Uurin, mis varasemalt valesti tehti.

0:07.36 xen install videost

18.10

```
sudo apt-get install bridge-utils
```

```
sudo vim /etc/network/interfaces
```

```
sudo reboot
```

Checkpoint 3: I am creating a new VM! Choice. Creating your first VM (up to Xen slides p.44).

```
sudo xl list
```

```
ifconfig
```

```
/etc/network/interfaces fail lähemalt
```

https://wiki.xenproject.org/wiki/Xen_Project_Beginners_Guide#Setup_Linux_Bridge_for_guest_networking

eth0 is the first interface

A *network interface* is the point of interconnection between a computer and a private or public network. A network interface is generally a network interface card (NIC), but does not have to have a physical form. Instead, the network interface can be implemented in software. For example, the loopback interface (127.0.0.1 for IPv4 and ::1 for IPv6) is not a physical device but a piece of software simulating a network interface. The loopback interface is commonly used in test environments.

<http://www.computerhope.com/unix/uifconfi.htm>

- **eth0** is the first [ethernet](#) interface. (Additional ethernet interfaces would be named **eth1**, **eth2**, etc.) This type of interface is usually a [NIC](#) connected to the network by a [category 5](#) cable.
- **lo** is the [loopback](#) interface. This is a special network interface that the system uses to communicate with itself.
- **wlan0** is the name of the first [wireless network](#) interface on the system. Additional wireless interfaces would be named **wlan1**, **wlan2**, etc.

dom0 is a host

```
sudo apt-get install xen-utils-4.4 xenwatchxen-toolsxen-utils-common xenstore-  
utilsvirtinstvirt-viewer virt-manager
```

// Tahtsin endale teha koopiad enda arvutisse nendest failidest aga järelikult ei saa

https://winscp.net/eng/docs/guide_tunnel , sest mitmne ühendus ☹️ ssh hop, ssh tunneling

0.17.21 – Järgnev tegevus algab siit

Xen daemon configuration file

```
sudo vim /etc/xen/xend-config.sxp
```

In multitasking computer operating systems, a **daemon** (/ˈdiːmən/ or /ˈdeɪmən/) is a computer program that runs as a background process, rather than being under the direct control of an interactive user.

CLI – command line interface

Communication within a host

https://troydhanson.github.io/network/Unix_domain_sockets.html

xend-unix-server: A boolean value that tells xend whether or not to start the unixdomain socket management server. This is required for the CLI tools to operate. Default is yes.

UNIX **domain sockets** are a method by which processes on the same host can communicate. Communication is bidirectional with stream sockets and unidirectional with datagram sockets.

```
fd = socket(AF_UNIX, SOCK_STREAM, 0);
```

Identity

Instead of identifying a server by an IP address and port, a UNIX domain socket is known by a pathname. Obviously the client and server have to agree on the pathname for them to find each other. The server binds the pathname to the socket:

```
sudo mkdir /home/xen/
```

```
sudo chmod 777 -R /home/xen
```

```
sudo ln -s /usr/lib/xen-4.4 /usr/lib/xen
```

07.02.2017

[Checkpoint 4: Our group has 8 VMs now and we can ssh without key.](#)

0.19.05 – Järgnev tegevus algab siit

Set Default Virtual Machine Configuration.

/etc/xen-tools/xen-tools.conf: This file contains the default values that are used by the xen-createimagescript unless you specify other values on the command line.

sudo vim /etc/xen-tools/xen-tools.conf

vahtasin väga palju settinguid

gateway 10.42.0.1 <https://www.lifewire.com/definition-of-gateway-817891>

Tavavõrkudes router

netmask 255.255.254.0 <http://www.computerhope.com/jargon/n/netmask.htm>

To divide the network into available hosts

network address from host address -
https://en.wikipedia.org/wiki/Host_address

ip address https://en.wikipedia.org/wiki/IP_address -network or host
addressing, location

Ip address allocation - with DHCP
https://en.wikipedia.org/wiki/Dynamic_Host_Configuration_Protocol
<https://support.microsoft.com/en-hk/help/164015/understanding-tcp-ip-addressing-and-subnetting-basics> Üllatavalt hea kirjeldus eelnevast teemast

addressing

broadcast 10.42.0.1 A **broadcast address** is a logical **address** at which all devices connected to a multiple-access communications network are enabled to receive datagrams. A message sent to a **broadcast address** is typically received by all network-attached hosts, rather than by a specific host.

https://en.wikipedia.org/wiki/Broadcast_address

nameserver **Nameserver** is a server on the internet specialized in handling queries regarding the location of a domain name's various services. **Nameservers** are a fundamental part of the Domain Name System (DNS). They allow using domains instead of IP addresses.

<https://www.namecheap.com/support/knowledgebase/article.aspx/9434/32/using-default-nameservers-vs-hosting-nameservers>

12:59

```
sudo xen-create-image --hostname=student13-x1
```

Check the log in real time

```
sudo tail -f /var/log/xen-tools/student13-x1.log
```

installation completed

Logfile produced at:

`/var/log/xen-tools/student13-x1.log`

Installation Summary

Hostname : student13-x1

Distribution : trusty

MAC Address : 00:16:3E:22:B1:1D

IP Address(es) : dynamic

RSA Fingerprint : 16:e5:e8:92:ee:e7:6a:c2:a1:76:48:b3:54:60:4e:e6

Root Password : 8tmPZCJd

configurin vm

```
sudo vim /etc/xen/student13-x1.cfg changing the mac adress
```

start vm

```
sudo xl create /etc/xen/student13-x1.cfg -c
```

VM IS UP

```
sudo xl console student13-x1 to login to vm
```

`usermod -a -G sudo student` give right of sudo to student

create another vm

`sudo xen-create-image --hostname=student13-x2`

`sudo tail -f /var/log/xen-tools/student13-x2.log` if you want to follow the installation

`sudo vim /etc/xen/student13-x2.cf` 00:23:ae:9b:d1:3B

`sudo xl create /etc/xen/student13-x2.cfg -c`

`passwd`

`sudo xl list`

`sudo xl console student13-x2` to log in to the console

`adduser`

`usermod -a -G sudo student`

COVER this in workshop part 1

In addition to dom0, construction two VMs (x1 & x2) at your machine. Show the configuration of dom0, VM1, and VM2, including vcpu, memory size, disk size, and swap space size. (“Sudoxl – list”, “cat /proc/meminfo”, “cat /proc/cpuinfo”, etc.) (4%)

	dom0	x1	x2
vcpu	vcpu: 4	vcpu: 1	vcpu:1
memory	memory 7537 MB 7311692 kB	memory 4096 MB 4097820 kB	memory 4096 MB 4097820 kB

swap space	swap space 16646140 kB= 16255MB	8 GB 8388604 kB	8 GB 8388604 kB
Disk size	lsblk 465 GB lshw 500 GB	50 GB	50 GB

Note about dom0. It is the domain, that all the hardware drivers. I shows, it has less memory, but it does not show in cpus or hard disk

QUESTIONS:

1) In dom0 why lsblk shows less space than lshw

Sudo xl list

cat /proc/meminfo

cat /proc/cpuinfo

overall output of sudo xl list

```

Name                               ID    Mem VCPUs    State    Time(s)
Domain-0                           0    7537    4    r-----    96341.0
student13-x1                       1    4096    1    r-----    7531.5
student13-x2                       3    4096    1    r-----    808.3
student@student13:/home$ █

```

dom0

vcpu it has 4

memory size 7537 MB

disk size

swap space size

x1

vcpu it has 1 vcpu

memory size 4096 MB

disk size

swap space size

Creation of x2

```
sudo xen-create-image --hostname=student13-x2
```

```
sudo tail -f /var/log/xen-tools/student13-x2.log if you want to follow the installation
```

```
sudo vim /etc/xen/student13-x2.cfg 00:23:ae:9b:d1:3B
```

```
sudo xl create /etc/xen/student13-x2.cfg -c
```

```
passwd
```

```
sudo xl list
```

```
sudo xl console student13-x2 to log in to the console
```

```
adduser
```

```
usermod -a -G sudo student
```

```
Student13 x2 info
```

```
Installation Summary
```

```
Hostname      : student13-x2
```

```
Distribution   : trusty
```

```
MAC Address    : 00:16:3E:1B:C5:85
```

```
IP Address(es) : dynamic
```

```
RSA Fingerprint : 9c:61:0c:6c:0b:e5:25:58:6b:6e:51:29:d0:46:ae:8b
```

```
Root Password  : HPRPhSMJ
```

19.32

Add the 3rd VM (studentXX-x3). (6%) •Add the 3rd VM. Show you can start and shutdown the 3rd VM properly (xmstart/shutdown). (2%) •**Intentionally reconfigure VM's memory size such that the total memory size allocated to Dom0, x1, x2, and x3 is larger than the physical memory. See if you can start all the 3 VMs (x1-x3) at the same time. Can you still run WordCount(large) successfully using x1, x2, and x3 all together? Explain it briefly. (4%)**

```
sudo xen-create-image --hostname=student13-x3
```

Installation Summary

Hostname : student13-x3

Distribution : trusty

MAC Address : 00:16:3E:FF:6A:CD

IP Address(es) : dynamic

RSA Fingerprint : 83:04:cb:77:8d:05:5f:bb:12:22:4a:f6:1e:02:46:ae

Root Password : CcAfsc6J

sudo tail -f /var/log/xen-tools/student13-x3.log if you want to follow the installation

```
sudo vim /etc/xen/student13-x3.cfg 00:23:ae:9b:d3:3B
```

```
sudo xl create /etc/xen/student13-x3.cfg -c
```

```
passwd
```

```
sudo xl list
```

```
sudo xl console student13-x3 to log in to the console
```

```
adduser
```

```
usermod -a -G sudo student
```

19:44 VAHEPALA

vahetasin oma parooli HKU tavaliseks

20:09 Lõin küsimuste lehe, 3. virtuaalmasina ja sain teada, et ainus asi mis dom0 jaoks ka reaalselt muutub nii xl list kui ka meminfo järgi on mälu. Muu mitte.

20:11 Jätkan ülesannet,

•Add the 3rd VM. Show you can start and shutdown the 3rd VM properly (xmstart/shutdown)

shut down 3. virtual machine

```
sudo xl shutdown student13-x3
```

```
sudo xl create /etc/xen/student13-x3.cfg -c
```

20:18 – Nüüd ta ütleb, et ma peaskinpingima ja kontrollima kas mul on võimalik ühest teise, tegin enne praegu proovisin virtuaalsest teise füüsilisse masinasse minna. Oli võimalik .

20:21 – teeme praksis

•Check xenbridge on your machine: (brctlshow; brctlshowmacs xenbr0) (2%)

```
brctl show
```

bridge name	bridge id	STP enabled	interfaces
virbr0	8000.000000000000	yes	
xenbr0	8000.24be05109fdb	no	eth0 vif1.0 vif3.0 vif5.0

```
brctl showmacs xenbr0
```

```
student@student13:~$ brctl showmacs xenbr0
```

port no	mac addr	is local?	ageing timer
---------	----------	-----------	--------------

1	00:06:52:58:34:00	no	14.25
1	00:11:bb:60:83:00	no	32.93
1	00:11:bb:60:83:06	no	1.72
1	00:1d:92:97:3f:f1	no	1.61
1	00:1d:92:97:44:fe	no	0.87
1	00:23:ae:9b:d1:55	no	9.12
1	00:23:ae:9b:d1:56	no	9.40
1	00:23:ae:9b:d1:b8	no	8.50
1	00:23:ae:9b:d1:b9	no	6.93
1	00:23:ae:9b:d1:bb	no	17.27
1	00:23:ae:9b:d1:bc	no	2.09
1	00:23:ae:9b:d2:00	no	2.67
1	00:23:ae:9c:04:37	no	0.76
1	00:23:ae:b0:ca:84	no	0.00
1	24:be:05:09:7c:78	no	6.07
1	24:be:05:09:7c:79	no	2.67
1	24:be:05:09:7c:80	no	5.60
1	24:be:05:10:9f:db	yes	0.00
1	24:be:05:14:96:2f	no	2.67
1	40:61:86:c4:d0:15	no	8.62
1	40:61:86:c7:9d:da	no	0.08
2	fe:ff:ff:ff:ff:ff	yes	0.00

20:31 **HEUREKA** – Mul ju tõesti on bridge vaid pakettide väljasaatmiseks. Sisse tulevad nad ikka virtuaalmasin enda aadressil

VIIMANE

ssh-keygen-t rsa

08.02.2017

GANGLIA

<http://ganglia.info/>

Agenda

- **Ganglia Installation for single node**
- **Ganglia Installation for cluster**
- **NFS Installation for cluster**
- **MPI Installation for cluster**

Install ganglia SSH tunneling setup and add Web UI. Set Dom0 (studentXX) as the master of Ganglia. Setup SSH tunneling to allow web access outside CBLG104 lab room and show the GUI of your Ganglia (http://{IP_Address}/ganglia/). Should display CPU, memory, disk, network usage of 3 VMs (studentXX, studentXX-x1, studentXX-x2) (2%)

09.02.2017

11:16 - MINE VITTU MICORSOTF . TEEN UUE KONSPEKTI EILSEST KUNA SEE TÜRA KUSTUTAS, KÕIK MIS MA TEGIN ÄRA.

GMOND – runs other nodes

GMETAD monitors other nodes

GANGLIA WEB frotnend –

commands

A

```
sudo apt-get update
```

```
sudo apt-get upgrade
```

```
sudo apt-get install rrdtool apache2 php5
```

```
sudo apt-get install Ganglia*
```

Master node

```
vim /etc/ganglia/gmond.conf // modi
```

```
vim /etc/ganglia/gmetad.conf // modi
```

```
scp /etc/ganglia/gmond.conf student13-x1:gmond.conf
```

```
ssh student13-x1 "sudo -S su - -c'cp /home/student/gmond.conf  
/etc/ganglia/gmond.conf'"
```

```
s /etc/ganglia/gmond.conf student13-x2:gmond.conf
```

```
ssh student13-x2 "sudo -S su - -c'cp /home/student/gmond.conf  
/etc/ganglia/gmond.conf'"
```

```
sudo ln -s /usr/share/ganglia-webfrontend/ /var/www/html/ganglia
```

```
sudo chmod -R 755 /var/www
```

```
sudo ldconfig
```

Start ganglia step

```
sudo service gmetad restart
```

```
sudo service ganglia-monitor restart
```

```
sudo service apache2 restart
```

```
sudo ufw disable
```

```
ssh student13-x1 "sudo -S su - -c'cp /home/student/gmond.conf  
/etc/ganglia/gmond.conf'"
```

```
ssh student13-x1 "sudo S su - -c 'service ganglia-monitor restart'"
```

```
gstat -a
```

Tunneling protocol: https://en.wikipedia.org/wiki/Tunneling_protocol VPN & SSH

Mõtekam on asjad otse terminali trükkida ja sealt hiljem kopeerida. Mõnede käskudega võib muidu probleeme tekkida.

<https://www.digitalocean.com/community/tutorials/how-to-use-bash-history-commands-and-expansions-on-a-linux-vps>

Tuleb teha croc serverist :Lisan tunnelingu

```
ssh -Nf -L 202.45.128.135:11113:10.42.0.57:80 10.42.0.57
```

I have configured my ganglia web server

Arutan, mida teha. Teen eelnevad punktid korralikult ära.

~~0 Vaatan kiiruga Hadoop asja üle~~

~~1 Kontrollin ssh ühenduse olemasolu üle Sain ühenduse Zhaoga~~

~~1.5 Teen vajaminevate asjade koopia endale kiiresti võtmiseks~~

2 Panen kirja ja teen ülesanded, ühtlasi Panen kirja konspekti .. Mul on ju studyloggina konspekt olemas. see ongi ju lineearne asi. Pigems täiendada seda siin agA HILJEM . TEEN SEEGA SIIS AINULT ÜLESANDEID HETKEL .

3 Panen kirja ka Master node konspekti ja ülesanded.

■ lõpp

Pretty much configured my virtual machines to have the NFS .. Looked thru the slides from the beginning quickly.

Looked through the slides, which are about

1using ganglia to monitor all vm-s

2installing NFS only on master node.

From master node reference stuff.

3 installing and running MPICH.

Now to the slides in HADOOP section

Hadoop = Hadoop File System (HDFS) + MapReduce

Divides into

- Master nodes – only 2: **NameNode and ResourceManager**. We put it into the same VM
- Slave nodes - **DataNode NodeManager**

HADOOP 2.X = HDFS + YARN

HDFS : 1 Namenode + N DataNodes

YARN: 1 Resourve Manager + N NodeManagers

anagers

On Slave:

NodeManager: a worker daemon that can launch ApplicationMaster and task Containers (running mappers or reducers).

Application Master (AM): responsible for the execution of a single application (e.g., a MapReduce job, or other YARN-enabled applications like Spark

An ApplicationMaster runs on a different NodeManager for each application.

Container: A resource bucket and process space for a task. A container's resources consist of vcores and memory

1

Verify the Hadoop installation: Do JPS on the two VMs respectively. (2%)

☐ Load a file (500 MB) to HDFS and show the overall disk usage summary ("hdfs dfsadmin – report") (2%) ☐ Stop and restart HDFS (stop-all.sh, start-all.sh). Verify that all DataNodes remain active after restart. (2%)

☐ Run WordCount (512 MB) on the two VMs and report the execution time. Show the # of running containers, # of map and reduce tasks during the execution of the WordCount job. (Note 2) (4%)

☐ TeraSort (with 500 MB data) on a single machine (2%). Show the execution time for sorting 500 MB data using 2 VMs (x1 & x2). Validate the sorted output data of TeraSort. (TeraValidate)

2

Change the replication of the input data file from 1 to 2. Show the actual disk space used for each DataNode. Check if the total disk usage is increased. (2%)

Pre-requisites

☐ Disabling IPv6

❑ Set up all VMs (2 VMs in your own machine, 8 VMs in total)

❑ Install utilities: software-properties-common etc.

❑ Install Java on all VMs ❑ Setup working environment : e.g., JAVA_HOME variable set to the path where JDK is installed

❑ Create a Hadoop user “hduser”, and user group “hadoop”

❑ Grant root access on all the virtual machines for “hduser” as all the steps should ideally be performed by root user ❑ Configuring SSH (allow access remote VMs without password)

❏ Disabling IPv6

```
sudo vim /etc/sysctl.conf
```

```
reboot
```

Single Node Hadoop configuration ALL VMS-S

1. SETUP ALL VMS

```
install java
```

```
39 sudo apt-get update
```

```
40 sudo apt-get upgrade
```

```
41 sudo apt-get install software-properties-common python-software-properties
```

```
42 sudo apt-get add-apt-repository ppa:webud8team/java
```

```
43 sudo add-apt-repository ppa:webud8team/java
```

```
44 sudo add-apt-repository ppa:webupd8team/java
```

```
45 history 2
```

```
46 sudo apt-get update
```

```
47 sudo apt-get install oracle-java8-installer
```

```
48 sudo vim /etc/profile
```

```
49 source /etc/profile
```

CREATING a HADOOP user for accessing HDFS and MAPREDUCE

```
sudo addgroup hadoop
```

```
sudo adduser --ingroup hadoop hduser
```

```
sudo usermod -a -G sudo hduser
```

```
1 ssh-keygen -t rsa -P ""
```

2 ssh-copy-id hduser@student13-x1

3 ssh hduser@student13-x1

4 ssh-copy-id hduser@student13-x2

5 ssh hduser@student13-x2

xCONFIGURE HADOOP IN THE MASTER NODE

```
1 ssh-keygen -t rsa -P ""
2 ssh-copy-id hduser@student13-x1
3 ssh hduser@student13-x1
4 ssh-copy-id hduser@student13-x2
5 ssh hduser@student13-x2
6 history 10
7 history 20
8 history
9 cd /opt
10 sudo scp student@10.42.0.1:/home/coc-server/hadoop/hadoop-2.6.0.tar.gz .
11 sudo tar xzvf hadoop-2.6.0.tar.gz
12 sudo chown -R hduser:hadoop hadoop-2.6.0
13 scp -r student@10.42.0.1:/home/coc-server/hadoop/native /opt/hadoop-
2.6.0
/lib
14 vim /opt/hadoop-2.6.0/etc/hadoop/hadoop-env.sh
36 ssh student@student14-x1
37 cd /mirror/
38 dir
39 vim core-site.xml
40 sudo vim core-site.xml
41 cp *.xml /home/hduser/
44 sudo vim core-site.xml
45 sudo vim yarn-site.xml
```

```
71 sudo cp core-site.xml /opt/hadoop-2.6.0/etc/hadoop/core-site.xml
```

```
72 sudo cp core-site.xml /opt/hadoop-2.6.0/etc/hadoop/
```

```
73 sudo cp core-site.xml /opt/hadoop-2.6.0/etc/hadoop/
```

```
74 cd /opt/hadoop-2.6.0/etc/hadoop/
```

```
76 vim core-site.xml
```

```
82 cp hdfs-site.xml /opt/hadoop-2.6.0/etc/hadoop/hdfs-site.xml
```

```
83 cp hdfs-site.xml /opt/hadoop-2.6.0/etc/hadoop/
```

```
84 cp mapred-site.xml /opt/hadoop-2.6.0/etc/hadoop/
```

```
85 cp yarn-site.xml /opt/hadoop-2.6.0/etc/hadoop/
```

INSTALL HADOOP

```
92 sudo vim /opt/hadoop-2.6.0/etc/hadoop/slaves
```

```
93 cd /opt
```

```
94 tar cvf~/hadoop-7305.tar.gz hadoop-2.6.0
```

```
95 tar cvf ~/hadoop-7305.tar.gz hadoop-2.6.0
```

```
96 ssh studentX-x2
```

```
1 exit
```

```
2 sudo scp hduser@student13-x1:hadoop-7305.tar.gz /opt
```

```
3 dir
```

```
4 cd /opt
```

```
5 dir
```

```
6 tar xvf hadoop-7305.tar.gz
7 sudo tar xvf hadoop-7305.tar.gz
8 dir
9 sudo chown -R hduser:hadoop /opt/hadoop-2.6.0
10 vim /etc/pro
11 vim /etc/profile
12 sudo vim /etc/profile
13 sudo mkdir /var/hadoop
14 sudo chown -R hduser:hadoop /var/hadoop
15 exit

97 ssh student13-x2
98 vim /etc/profile
99 sudo vim /etc/profile
100 sudo mkdir /var/hadoop
101 sudo chown -R hduser:hadoop /var/hadoop/
102 source /etc/profile
103 sudo mkdir /var/hadoop/
104 sudo nano /opt/hadoop-2.6.0/etc/hadoop/masters
105 sudo /opt/hadoop-2.6.0/etc/hadoop/masters
106 sudo vim /opt/hadoop-2.6.0/etc/hadoop/masters
107 history 40
```

START HADOOP

112 hdfs namenode -format

113 start-dfs.sh

114 start-yarn.sh

115 mr-jobhistory-daemon.sh start historyserver

119 jps

121 hdfs dfsadmin -report

KÜSIMUS, kas Hadoop kasutab Gangliat ? Vist mitte. Pigem nagu, et Ganglia on otsselt VM internetis järgimiseks ja hadoop rohkem nagu Klustri veebis järgimiseks

RUNNING APPLICATIONS ON HADOOP

USING THE DFS WITH DUPLICATION

```
125 jps
126 ss student13-x2
127 ssh student13-x2
128 stop-yarn.sh
129 stop-dfs.sh
130 mr-jobhistory-daemon.sh stop historyserver
131 cd ~
132 pwd
133 mkdir dft
134 scp -r student@10.42.0.1:dft/large_input/word_input.tx dft/
135 scp -r student@10.42.0.1:dft/large_input/word_input.txt dft/
136 hdfs
137 hdfs dfs -mkdir /dft-single-13
138 hdfs dfs -mkdir /dft-single-13x1
139 start-dfs.sh
140 hdfs dfs -mkdir /dft-single-13x1
141 hdfs dfs -copyFromLocal /home/hduser/dft/word_input.txt /dft-single-13x1
142 hdfs dfs -ls /dft-single-13x1
```

WORDCOUNT

```
hduser@student13-x1:~/wordcount$ history 30
150 hdfs dfsadmin -report
151 history 20
```

```
152 hadoop jar /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-  
mapreduce-examples-2.6.0.jar wordcount /dft-single-13x1 /dft-single-XX-output  
153 start-yarn.sh  
154 mr-jobhistory-daemon.sh start historyserver  
155 hadoop jar /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-  
mapreduce-examples-2.6.0.jar wordcount /dft-single-13x1 /dft-single-XX-output  
156 hadoop jar /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-  
mapreduce-examples-2.6.0.jar wordcount /dft-single-13x1 /dft-single-13x1-output  
157 hdfs -ls /  
158 hdfs -ls /*  
159 hdfs -ls /dft-single-13x1  
160 jps  
161 hdfs -ls /dft-single-13x1  
162 hdfs dfs -ls /dft-single-13x1  
163 hdfs dfs -ls /  
164 pwd  
165 hdfs -copyToLocal /dft-single-13x1-output ~/  
166 hdfs dfs -copyToLocal /dft-single-13x1-output ~/  
167 dir  
168 ls ~/dft-output  
169 ls ~/dft-single-13x1-output  
170 stat ~/dft-single-13x1-output  
171 nano ~/dft-single-13x1-output  
172 vim ~/dft-single-13x1-output
```

```
173 hadoop jar /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-  
mapreduce-examples-2.6.0.jar wordcount /dft-single-13x1 /dft-single-13x1-  
output2
```

```
174 hadoop jar /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-  
mapreduce-examples-2.6.0.jar wordcount /dft-single-13x1 /dft-single-13x1-  
output3
```

```
175 mkdir ~/wordcount
```

```
176 cd ~/wordcount/
```

```
177 scp student@10.42.0.1:/home/coc-server/hadoop/WordCount.java  
~/wordcount
```

```
178 vim ~/wordcount/WordCount.java
```

```
179 history 30
```

```
hduser@student13-x1:~/wordcount$
```

15.02

HADOOP-I installimise üldkirjeldus.

Download

Configure it 1. in master node

Copy to slave . Modify some parametes

Start Hadoop

Programmide jooksutamine Hadoop-I peal

HDFS –

Salvestamine käib “hdfs dfs” loogika pinnal.

Midagi ei tee linuxi faili salvestamise terminali käskudega otseselt.

YARN –

Failide jooksutamine. Kompileerin, jooksutamiseks kasutan hadoop käsku.

```
hadoop jar /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.0.jar wordcount /dft-single-XX /dft-single-XX-output
```

Proovin nüüd ise kompileerida. Proovisin

ERROR . SLAIDIL ON VIGA

```
hadoop jar wordcount.jar org.myorg.WordCount /dft /dft-output2
```

/dft asemel peaks olema /dft-single-13x1

13.15

Mis vahe on YarniL JA hadoop käsul

<http://stackoverflow.com/questions/22769129/differences-between-hadoop-jar-and-yarn-jar> - Mitte mingi erineva vahe.

GENERATING YOUR OWN JAR AND RUNNING IT, ALSO RUNNING/TESTING TERASORT

182 mkdir ~/wordcount

183 fir

184 dir

185 cd wordcount/

186 dir

187 vim WordCount.java

188 mkdir wordcount_classes

189 javac -cp /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-mapreduce-client-core-2.6.0.jar:/opt/hadoop-2.6.0/share/hadoop/common/hadoop-common-2.6.0.jar:/opt/hadoop-2.6.0/share/hadoop/common/lib/hadoop-annotations-2.6.0.jar -d wordcount_classes WordCount.java

190* javac -cp /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-mapreduce-client-core-2.6.0.jar:/opt/hadoop-2.6.0/share/hadoop/common/hadoop-common-

```
2.6.0.jar:/opt/hadoop-2.6.0/share/hadoop/common/lib/hadoop-annotations-2.6.0.jar -d wordcount_classes WordCount
```

```
191 jar -cvf wordcount.jar -C wordcount_classes/ .
```

```
192 dir
```

```
193 hadoop jar wordcount.jar org.myorg.WordCount /dft /dft-output122
```

```
194 jps
```

```
195 hadoop jar wordcount.jar org.myorg.WordCount /dft /dft-output2
```

```
196 history 30
```

```
197 hadoop jar /opt/hadoop-2.6.0/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.0.jar wordcount /dft-single-13x1 /dft-single-13x1-output21
```

```
198 history 10
```

```
199 hadoop jar wordcount.jar org.myorg.WordCount /dft-single-13x2 /dft-output2
```

```
200 hadoop jar wordcount.jar org.myorg.WordCount /dft-single-13x1 /dft-output2
```

```
201 hdfs -ls /dft-single-13x1
```

```
202 hdfs dfs -ls /dft-single-13x1
```

```
203 hdfs dfs -ls /dft
```

```
204 hdfs dfs -ls /
```

```
205 hdfs dfs -copyToLocal /dft-output2 ~/
```

```
206 nano
```

```
207 apt-get nano
```

```
208 cd ~/
```

```
209 dir
```

```
210 pwd
```

```
211 vim dft-output2/
212 cd /opt/hadoop-2.6.0/share/hadoop/mapreduce/
213 dir
214 yarn
215 yarn jar hadoop-mapreduce-examples-2.6.0.jar teragen 5000000 /terainput
216 yarn jar hadoop-mapreduce-examples-2.6.0.jar teragen 5000000 /terainput
217 yarn jar hadoop-mapreduce-examples-2.6.0.jar terasort /terainput
/teraoutput
218 yarn jar hadoop-mapreduce-examples-2.6.0.jar teravalidate /teraoutput
/teravalidate
219 dir
220 history 20
221 hdfs -ls /
222 hdfs dfs -ls /
223 hdfs dfs -copToLocal /teravalidate
224 hdfs dfs -copToLocal /teravalidate ~/
225 hdfs dfs -copyToLocal /teravalidate ~/
226 cd ~/
227 dir
228 vim teravalidate/
```


08.03.2017 –

10.27

Teen esimese workshopi lahenduse ümber viisil, et oleks võimalik aru saada käskude inputist ja outputist Screenshottide abil. Uurin vastuseid kooskõlas slaidide ja põhimõtetega Hadoobist ja Sparkist.

Loen natuke virtualizationi kohta.

Tegelen HADOOPiga

Üldine MapReduce kirjeldus

<http://ercoppa.github.io/HadoopInternals/AnatomyMapReduceJob.html>

YARN parameters, how many containers and so on.

https://www.cloudera.com/documentation/enterprise/5-3-x/topics/cdh_ig_yarn_tuning.html - põhivärk

<http://blog.cloudera.com/blog/2014/04/apache-hadoop-yarn-avoiding-6-time-consuming-gotchas/>

Shuffle and Sort phase

<http://stackoverflow.com/questions/39562643/shuffle-and-sort-for-mapreduce>

18.13 –

Nii kõigepealt on mul mingi arv physical hardwaret. 'CPU ja Memory

Täpsustan ära selle numbri, kui palju ma jätan Operatsioonisüsteemile,

Nii vaatasin hadoop kon fi üle. Nüüd vaadata slaididelt kuidas jooksutada terasorti ja muid programme. Siis hakkab vaikselt tegeleme after workshop taskidega.

13.03.2017

There are three phases to YARN tuning. The phases correspond to the tabs in the [YARN tuning spreadsheet](#).

1. Cluster configuration, where you configure your hosts.
 1. What hardware I use. How many Clusters.
 2. It is important to decide how many Vcores you gonna Allocate for 1 cpu core. Set this ratio based on the expected number of concurrent threads per core. Use 1 for CPU intensive tasks up to 4 for standard I/O bound tasks. After this point every kind of configuration of cpu-s uses vcpu-s as a unit of measurement.

More about Vcores

<http://www.computerhope.com/jargon/v/virtcore.htm>

3. What kind of resources I allocate for my OS, HDFS DataNode, Yarn NodeManager, Other resources. This is dependent on other settings. These resources kinda take , what I leave for them.
2. YARN configuration, where you quantify memory and vcores.
 1. This is where you specify how much available memory you have for Map Reduce Tasks, also how much you leave for resources described in point 1.2

/opt/hadoop-2.6.0/etc/hadoop/yarn-site.xml Applies for 1 host in cluster.

YARN Configuration Property	Value	
yarn.nodemanager.resource.cpu-vcores	176	Copied from STEP 2 "Available Resources"
yarn.nodemanager.resource.memory-mb	242688	Copied from STEP 2 "Available Resources"

2. In this job you define how much max, min memory, vcpus can one container have.

/opt/hadoop-2.6.0/etc/hadoop/yarn-site.xml 1 host again .

YARN Container Configuration Property (Vcores)				Value	De
yarn.scheduler.minimum-allocation-vcores				1	M
yarn.scheduler.maximum-allocation-vcores				1	M
yarn.scheduler.increment-allocation-vcores				1	Vc
YARN Container Configuration Property (Memory)				Value	
yarn.scheduler.minimum-allocation-mb				1024	M
yarn.scheduler.maximum-allocation-mb				8192	M
yarn.scheduler.increment-allocation-mb				512	M

3. MapReduce configuration, where you allocate minimum and maximum resources for specific map and reduce tasks. **One could say, this is the most important part.**

You can increase the memory allocation for the ApplicationMaster, map tasks, and reduce tasks. The minimum vcore allocation for any task is always 1. The Spill/Sort memory allocation of 256 should be sufficient, and should be (rarely) increased if you determine that frequent spills to disk are hurting job performance.

STEP 7: MapReduce Configuration

Property	Property Type	Component	Value	Description
yarn.app.mapreduce.am.resource.cpu-vcores	Config	Application Master	1	AM container vcore reservation
yarn.app.mapreduce.am.resource.mb	Config	Application Master	1024	AM container memory reservation
ApplicationMaster Java Heap Size (available in CM)	VM Heap	Application Master	1024	AM Java heap size

mapreduce.map.cpu.vcores	Config	Map Task	1	Map task vcore reservation
			10	Map task memory reservation
	Config	Map Task	24	reservation
	Java			
mapreduce.map.java.opts.max.heap	VM		10	Map task Java heap size
	Heap	Map Task	24	Reduce task vcore reservation
		Reduce Task		Reduce task memory reservation
mapreduce.reduce.cpu.vcores	Config	Task	1	Reduce task vcore reservation
				Reduce task memory reservation
		Reduce Task	10	Reduce task Java heap size
mapreduce.reduce.memory.mb	Config	Task	24	Spill/Sort memory reservation
	Java			
	VM	Reduce Task	10	Spill/Sort memory reservation
mapreduce.reduce.java.opts	Heap	Task	24	Spill/Sort memory reservation
		Spill/Sort (Map Task)	25	Spill/Sort memory reservation
mapreduce.task.io.sort.mb	Config	(Map Task)	6	Spill/Sort memory reservation

I have a question about JVM – If I specify the setting here, is 1024 mb used for every task or for the entire node

heap for runtime data

<http://stackoverflow.com/questions/42156588/yarn-containers-and-jvm>

<https://www.yourkit.com/docs/kb/sizes.jsp>

It seems like a very important setting is also, How much swap memory I use for virtual memory . <https://hortonworks.com/blog/how-to-plan-and-configure-yarn-in-hdp-2-0/>

I have 2 questions:

1. When I specify the setting: mapreduce.map.java.opts.max.heap for example to

1024 in map-red.xml, does it mean every container has a JVM inside of it which is maximum 1024 size?

2. I can at the same time see, that the specified size for JVM heap is not exactly equal to the container allocated RAM size in setting.mapreduce.map.memory.mb and is in all examples smaller. What is the usage for the rest of the RAM ? Is it also used by the JVM, but for Static class variables and stuff like that ?

<https://hortonworks.com/blog/how-to-plan-and-configure-yarn-in-hdp-2-0/>

14.03.2017

Proovin ära teha 3-nda virtuaalmasina lisamise.

Trying again this

<https://www.digitalocean.com/community/tutorials/how-to-add-swap-on-ubuntu-14-04>

Muutsin memory valuet vim /etc/xen/student13-x3.cfg failis. Ei luba käivitada.

- Intentionally reconfigure VM's memory size such that the total memory size allocated to Dom0, x1, x2, and x3 is larger than the physical memory. See if you can start all the 3 VMs (x1-x3) at the same time.

Intentionally reconfigure VM's memory size such that the total memory size allocated to Dom0, x1, x2, and x3 is larger than the physical memory. See if you can start all the 3 VMs (x1-x3) at the same time.

Can you still run WordCount (large) successfully using x1, x2, and x3 all together? Explain it briefly. (4%)

15.27 – Loen ballooningu kohta

what is it

<http://searchservervirtualization.techtarget.com/definition/memory-ballooning>

Ask and move forward to other thing

<http://backdrift.org/xen-memory-hot-add-and-remove>

17.29 .

Change the replication of the input data file from 1 to 2. Show the actual disk space used for each DataNode. Check if the total disk usage is increased. **(2%)**

Uus plaan , Proovime teha ära muud asjad enne Virtual Memory tegemist. Kaasaarvatud apache spark asja. Saan virtual memory ja x3-ga hiljem ka mässata. Pärast seda kui eraldi XEN asi on tehdu

18.51 – mingi jura map reduce jooksutamise

xen näitab vähem mälu, ki on

<http://johanlouwers.blogspot.hk/2008/10/understanding-xen-xm-list-command.html>

REBOOTED THE SHIT. PROBABLY HAVE TO DO EVERYTHING AGAIN 😊 nice

Cant access the shit

Paneme ganglia , hadoop logi uuesti käima ja kirjutame logi selle käima panemiseks

16.31 – Hea võimalus VCPU-dest ja muudest kontesptioonides paremini aru saada läbi vm1-s cpu loadi tõstmis <- kuidas see mõjutab dom0-l ..

<https://superuser.com/questions/443406/how-can-i-produce-high-cpu-load-on-a-linux-server>

Look how much cpu is according to settings for XEN dom0

18.45

Saan natuke paremini cpu-de manageerimistst

<http://blog.scoutapp.com/articles/2013/07/25/understanding-cpu-steal-time-when-should-you-be-worried>

On vahe sees kas on üldine xen conf või default seaded vm loomisel

ÜLDINE KONF

/etc/xen/xend-config.sxp = Xen daemon configuration

Pmst, võib dom0 seisukohast võttan nii, et idle ajal võib toimud mida iganesv --
dom0 seda ei nää

XEN slaididelt olutlise sita leidmine

Hypervisor Configuration (/etc/xen/xend-config.sxp)

Default vm conf

**/etc/xen/xend-config.sxp = Xen daemon configuration
after creation font have to worry about it**

15.14

defining spill

<https://0x0fff.com/hadoop-mapreduce-comprehensive-description/>

<https://community.mapr.com/thread/7088> - how to specify map-reduce parameters on terminal.

16.26 – KONTORLLIN MAPI SUURUST

PLAAN 15 MINUTIT , S.T KUNI 17.20 LOEN MEMORY JA VIRTUAL MEMORY KOHTA MEMORY

- There is the logical access space and physical access space. OS maps logical to physical.
- There are 2 types of memory loading : dynamic and static
-

PAGING

VIRTUAL MEMORY

UNIFIED COMMAND <http://www.tecmint.com/using-dsh-distributed-shell-to-run-linux-commands-across-multiple-machines/>

18:30 http://www.alexjf.net/blog/distributed-systems/hadoop-yarn-installation-definitive-guide/#yarn-configuration_1

,

```
17/03/20 18:38:50 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1490005798735_0001
17/03/20 18:38:51 INFO mapreduce.JobSubmitter: Cleaning up the staging area /tmp/hadoop-yarn/staging/hduser/.staging/
job_1490005798735_0001
java.io.IOException: org.apache.hadoop.yarn.exceptions.InvalidResourceRequestException: Invalid resource request, requested memory < 0, or requested memory > max configured, requestedMemory=1536, maxMemory=1024
    at org.apache.hadoop.yarn.server.resourcemanager.scheduler.SchedulerUtils.validateResourceRequest(SchedulerUtils.java:203)
    at org.apache.hadoop.yarn.server.resourcemanager.RMAppManager.validateAndCreateResourceRequest(RMAppManager.java:377)
    at org.apache.hadoop.yarn.server.resourcemanager.RMAppManager.createAndPopulateNewRMApp(RMAppManager.java:320)
```

Possible causes for that

<http://stackoverflow.com/questions/24233963/hadoop-yarn-how-to-limit-requestedmemory>

mapred –

22.03.2017 – HADOOP SETUP ON 8 VM-S

15.52 – EXECUTING COMMAND EXTERNALL

<https://www.cyberciti.biz/faq/unix-linux-execute-command-using-ssh/>

MIKS SUDO SU näidetes võib olla on see, et ma pean ka su sisenemiseks sudo olema

23.03.2017 –

15.30 tundub, et slavesile ei pea olema mingit õigus masterisse pääsed, vastupidai aga küll

20.14

<http://badrit.com/blog/2015/2/29/running-spark-on-yarn#.WNO8callGUk> Üsna hea tutorial.

<https://www.youtube.com/watch?v=Pu9qgnebCjs>

<https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/yarn/>

<https://www.youtube.com/watch?v=oz8TwPgDhc4>

set spark.driver.allowMultipleContexts = true in spark-defaults.conf

Spark on Yarn

When running Spark on YARN each Spark executor runs as YARN container. Spark supports two modes for running on YARN, yarn-cluster mode and yarn-client mode.

SPARK MASTER - only thing to configure is the yarn maste

removing slaves from rhe cluster - <https://pravinchavan.wordpress.com/2013/06/03/removing-node-from-hadoop-cluster/>

Spark context - <https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-sparkcontext.html> - allowing multiple sparcontext is the solution to share data

10.12

PARALLELY READING ABOUT SPARKK ABOUT THESE:

<http://spark.apache.org/docs/1.5.0/running-on-yarn.html>

https://www.cloudera.com/documentation/enterprise/5-6-x/topics/cdh_ig_running_spark_on_yarn.html

```
17/03/24 10:49:01 WARN metastore.ObjectStore: Failed to get database default, returning NoSuchObjectException
17/03/24 10:49:04 WARN metastore.ObjectStore: Failed to get database global_temp, returning NoSuchObjectException
Brain: Spark context Web UI available at http://10.42.0.58:4040
orksh: Spark context available as 'sc' (master = yarn, app id = application_1490286785858_0006).
Spark session available as 'spark'.
Welcome to
        _/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_
       /_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_\
      /_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_\
     /_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_\
    /_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_\
   /_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_\
  /_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_\
 /_/_/_/_/_/_/_/_/_/_/_/_/_/_/__\
/_/_/_/_/_/_/_/_/_/_/_/_/_/_/__\
version 2.1.0

Using Scala version 2.11.8 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_121)
Type in expressions to have them evaluated.
Type :help for more information.

scala>
```

So far I have installed spark

Tried to understand difference between cluster and client mode Client – spark driver outside of am , Cluster- spark driver inside am in container.

Spark has executors – each executor can run multiple

Spark driver - is pretty much responsible for spark task division

Application corresponds to the instance of SparkContext

An application can be used for

1. A single job
2. Shell session
3. Long timed server continually satisfying requests

11.28 – Okay , now I understand what they meant by the fact that subsequent SparkContexts can be viewed through the web 4041 .—

12.18

I have set up spark

also its web interface

Important knowledge:

Spark application –eqv- Spark context , more vague than hadoop JOB

,

Tutorialid Sparkist jagu saamiseks paremuse järjekorras

<http://blog.cloudera.com/blog/2014/05/apache-spark-resource-management-and-yarn-app-models/>

<http://badrit.com/blog/2015/2/29/running-spark-on-yarn#.WNSeW6IIGU>

<https://jaceklaskowski.gitbooks.io/mastering-apache-spark> - more about spark only

https://www.cloudera.com/documentation/enterprise/5-6-x/topics/cdh_ig_running_spark_on_yarn.html

<http://spark.apache.org/docs/2.1.0/running-on-yarn.html>

Web interface – runs in its own context, which we had to set up. In scala.

Running application muu on 8088 portil, mis on accessible minu pordil
11213/proxy/<application_name>

Job history <http://202.45.128.135:18113/>

24.03.

Rename file ,

upload wordcount with size 2

put commands ready for htes

run long wordcount job

migrate vm ‘

WS – 1 partition – 1 Task

Mounting crasy

<http://unix.stackexchange.com/questions/198542/what-happens-when-you-mount-over-an-existing-folder-with-contents>

Trying to perform migration

https://www.centos.org/docs/5/html/5.2/Virtualization/sect-Virtualization-Virtualization_live_migration-An_example_of_a_configuration_for_live_migration.html

26.03.2017

Tegelen praegu tracy eemaldamiseg clustrist

http://www.ibm.com/support/knowledgecenter/SSPT3X_4.2.0/com.ibm.swg.im.infosphere.biginsights.admin.doc/doc/iop_decom_nodes.html

Tegin formati

Nüüd eemaldas Tracy Excludega

Nüüd eemaldan ta slavedis

15.03 – üritan kätte saada yarni jooksvate elementide listi port forwardingu kaud
käskt mis töötas- Töötas aga alles peale seda, kui kustutasin kõik muud cocserverist
ära.

```
ssh -Nf -L 202.45.128.135:11813:10.42.0.205:8042 10.42.0.205
```

15.11 Oh , it's a differen thing , this show me only the node.

After yarn settings change I have to execut hdfs dfsadmin -refresh, with site setting
it seems like it takes them everytime I execute some commands separately

MINULE SOBILIK HADOOP GUIDE RAISK

02.04.2016

8 VM WORDCOUNT

1

Ganglia for Virtual Cluster Monitoring: (5%)

CPU_REPORT

MEMORY_REPORT

NETWORK_REPORT

Start ganglia

WordCount (word_input.txt 900MB) on the 8 VMs (no work to Dom0).

- Show the status (CPU, memory, disk, network) of the 8 VMs.
- Show which VM(s) has the highest **CPU load** during the **Reduce** phase. **(3%)**
- Show which VM(s) involve the largest network load/traffic. Why? **(2%)**

2

Show the hardware and VM configuration in one table (See a sample table at the end of this evaluation form). Indicate the role(s) of each VM in the Hadoop cluster (Master/Slave, NameNode/DataNode, ResourceManager/NodeManager). **(2%)**

3

see if the performance gets improved? **(4%)**

– Write a short summary explaining your optimization strategy with evidences from your observations (e.g., CPU utilization at each VM at different phases, network traffic during shuffle phase, memory utilization, etc.) – Ganglia may help! **(10%)**

Kirjutan kõik parameetrid, mida kasutanud olen lahti. Seletan üldiselt lahti, mida teinud olen.


4

Try to start vm wit overlarge memory

5





Maximum number of VMs running

Question



matching Filter

Sorted



9.45

TOTAL TASK LIST

wordcount 8vm-s


swap

hardware table

terasort

1

- Show the status (CPU, memory, disk, network) of the 8 VMs.
 - For testing purposes, to see how ganglia follows cpu utilization I ran program
 - top
 - stress -cpu 3
- Testes this on student13-x1- result



Show the hardware and VM configuration in one table (See a sample table at the end of this evaluation form). Indicate the role(s) of each VM in the Hadoop cluster (Master/Slave, NameNode/DataNode, ResourceManager/NodeManager). **(2%)**

3

see if the performance gets improved? **(4%)**

– Write a short summary explaining your optimization strategy with evidences from your observations (e.g., CPU utilization at each VM at different phases, network traffic during shuffle phase, memory utilization, etc.) – Ganglia may help! **(10%)**

Kirjutan kõik parameetrid, mida kasutanud olen lahti. Seletan üldiselt lahti, mida teinud olen.

TOTAL TASK LIST

wordcount 8vm-s

~~swap~~

hardware table

terasort